

# Homework 5

PSTAT 131/231

## Contents

Elastic Net Tuning . . . . .	1
------------------------------	---

## Elastic Net Tuning

For this assignment, we will be working with the file "pokemon.csv", found in /data. The file is from Kaggle: <https://www.kaggle.com/abcsds/pokemon>.

The Pokémon franchise encompasses video games, TV shows, movies, books, and a card game. This data set was drawn from the video game series and contains statistics about 721 Pokémon, or “pocket monsters.” In Pokémon games, the user plays as a trainer who collects, trades, and battles Pokémon to (a) collect all the Pokémon and (b) become the champion Pokémon trainer.

Each Pokémon has a primary type (some even have secondary types). Based on their type, a Pokémon is strong against some types, and vulnerable to others. (Think rock, paper, scissors.) A Fire-type Pokémon, for example, is vulnerable to Water-type Pokémon, but strong against Grass-type.



Figure 1: Fig 1. Vulpix, a Fire-type fox Pokémon from Generation 1.

The goal of this assignment is to build a statistical learning model that can predict the **primary type** of a Pokémon based on its generation, legendary status, and six battle statistics.

Read in the file and familiarize yourself with the variables using `pokemon_codebook.txt`.

```
library(tidymodels)
library(tidyverse)
library(ISLR) # For the Smarket data set
library(ISLR2) # For the Bikeshare data set
library(discrim)
```

```
library(poissonreg)
library(corr)
library(klaR) # for naive bayes
library(forcats)
library(corrplot)
library(pROC)
library(recipes)
library(rsample)
library(parsnip)
library(workflows)
library(janitor)
library(glmnet)
tidymodels_prefer()
set.seed(2022)
```

## Exercise 1

Install and load the `janitor` package. Use its `clean_names()` function on the Pokémon data, and save the results to work with for the rest of the assignment. What happened to the data? Why do you think `clean_names()` is useful?

```
# install.packages("janitor")
pokemon_raw <- read.csv("Pokemon.csv")
head(pokemon_raw)
```

```
##   X.      Name Type.1 Type.2 Total HP Attack Defense Sp..Atk
## 1  1      Bulbasaur  Grass Poison   318 45    49    49    65
## 2  2      Ivysaur   Grass Poison   405 60    62    63    80
## 3  3      Venusaur  Grass Poison   525 80    82    83   100
## 4  3 VenusaurMega Venusaur  Grass Poison   625 80   100   123   122
## 5  4      Charmander   Fire        309 39    52    43    60
## 6  5      Charmeleon   Fire        405 58    64    58    80
##   Sp..Def Speed Generation Legendary
## 1      65   45           1      False
## 2      80   60           1      False
## 3     100   80           1      False
## 4     120   80           1      False
## 5      50   65           1      False
## 6      65   80           1      False
```

```
pokemon1 <- clean_names(pokemon_raw)
head(pokemon1)
```

```
##   x      name type_1 type_2 total hp attack defense sp_atk sp_def
## 1 1      Bulbasaur  Grass Poison   318 45    49    49    65    65
## 2 2      Ivysaur   Grass Poison   405 60    62    63    80    80
## 3 3      Venusaur  Grass Poison   525 80    82    83   100   100
## 4 3 VenusaurMega Venusaur  Grass Poison   625 80   100   123   122   120
## 5 4      Charmander   Fire        309 39    52    43    60    50
## 6 5      Charmeleon   Fire        405 58    64    58    80    65
##   speed generation legendary
```

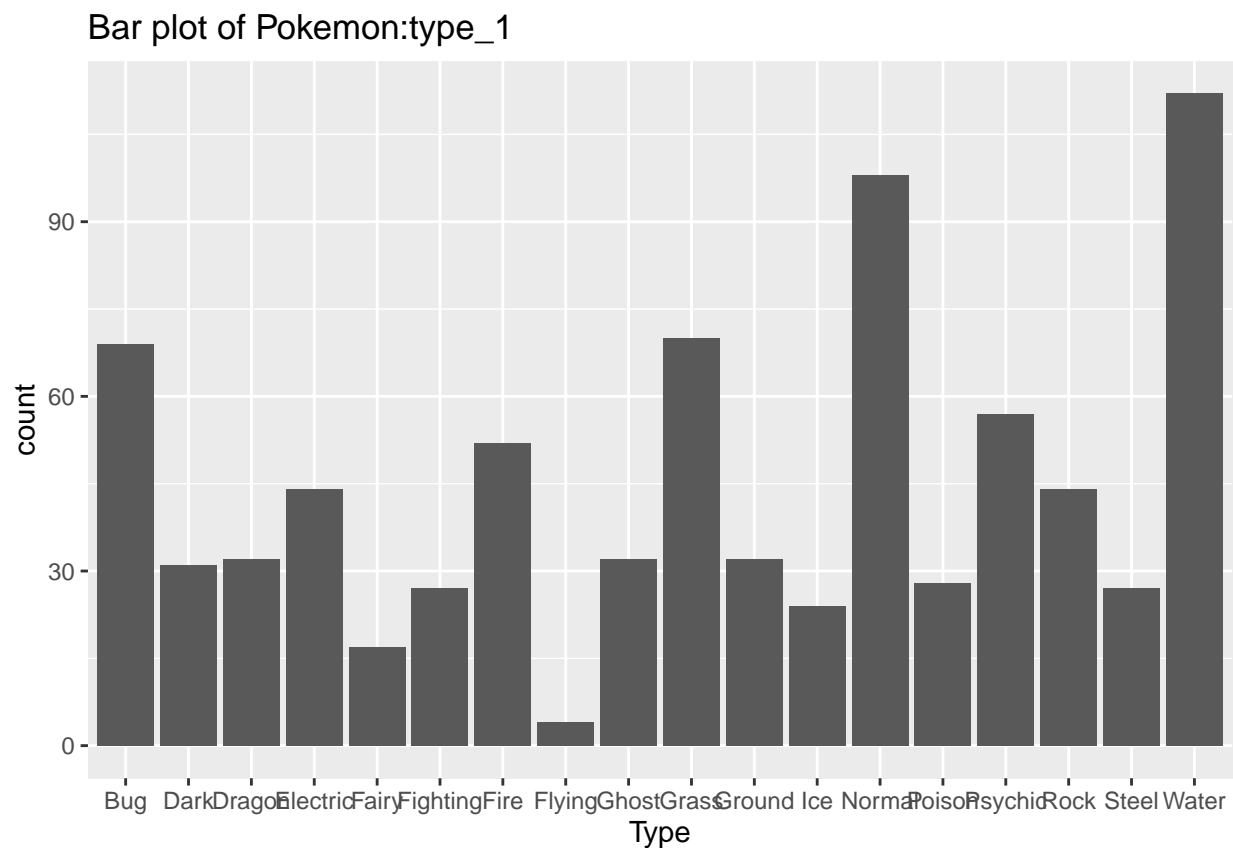
```
## 1    45      1    False
## 2    60      1    False
## 3    80      1    False
## 4    80      1    False
## 5    65      1    False
## 6    80      1    False
```

\*All column name convert to lower case and all of them are unique, also the name consist '?' convert to '\_.'  
 \*Resulting names are unique and consist only of the "\_" character, numbers, and letters which is more efficiency.\*

## Exercise 2

Using the entire data set, create a bar chart of the outcome variable, `type_1`.

```
ggplot(pokemon1, aes(x= type_1)) +
  geom_bar(stat = "count") +
  ggtitle("Bar plot of Pokemon:type_1") +
  xlab("Type")
```



How many classes of the outcome are there? Are there any Pokémon types with very few Pokémon? If so, which ones?

*Class: 18.*

*The flying type with few Pokemon.*

For this assignment, we'll handle the rarer classes by simply filtering them out. Filter the entire data set to contain only Pokémon whose `type_1` is Bug, Fire, Grass, Normal, Water, or Psychic.

After filtering, convert `type_1` and `legendary` to factors.

```
pokemon <- pokemon1[which(pokemon1$type_1 == "Bug" | pokemon1$type_1 == "Fire" |
  pokemon1$type_1 == "Grass" | pokemon1$type_1 == "Normal" |
  pokemon1$type_1 == "Water" | pokemon1$type_1 == "Psychic"), ]
pokemon <- pokemon %>%
  mutate(type_1 = factor(type_1),
         legendary = factor(legendary))
head(pokemon)
```

```
##   x          name type_1 type_2 total hp attack defense sp_atk sp_def
## 1 1      Bulbasaur Grass Poison  318 45    49    49    65    65
## 2 2      Ivysaur  Grass Poison  405 60    62    63    80    80
## 3 3      Venusaur Grass Poison  525 80    82    83   100   100
## 4 3 VenusaurMega Venusaur Grass Poison  625 80   100   123   122   120
## 5 4      Charmander  Fire    309 39    52    43    60    50
## 6 5      Charmeleon  Fire    405 58    64    58    80    65
##   speed generation legendary
## 1    45          1      False
## 2    60          1      False
## 3    80          1      False
## 4    80          1      False
## 5    65          1      False
## 6    80          1      False
```

### Exercise 3

Perform an initial split of the data. Stratify by the outcome variable. You can choose a proportion to use. Verify that your training and test sets have the desired number of observations.

```
pokemon_split <- pokemon %>%
  initial_split(strata = type_1, prop = 0.7)
pokemon_train <- training(pokemon_split)
pokemon_test <- testing(pokemon_split)
dim(pokemon_train)
```

```
## [1] 318 13
```

```
dim(pokemon_test)
```

```
## [1] 140 13
```

Next, use  $v$ -fold cross-validation on the training set. Use 5 folds. Stratify the folds by `type_1` as well. *Hint: Look for a `strata` argument.* Why might stratifying the folds be useful?

```
pokemon_fold <- vfold_cv(pokemon_train, v = 5, strata = type_1)
```

*Each resample is created within the stratification variable that each fold is an appropriate representative of the original data.*

## Exercise 4

Set up a recipe to predict `type_1` with `legendary`, `generation`, `sp_atk`, `attack`, `speed`, `defense`, `hp`, and `sp_def`.

- Dummy-code `legendary` and `generation`;
- Center and scale all predictors.

```
pokemon_recipe <- recipe(type_1 ~ legendary + generation + sp_atk +
                          attack + speed + defense + hp + sp_def, pokemon_train) %>%
  step_dummy(legendary, generation) %>%
  step_center(all_predictors()) %>%
  step_scale(all_predictors())
```

## Exercise 5

We'll be fitting and tuning an elastic net, tuning `penalty` and `mixture` (use `multinom_reg` with the `glmnet` engine).

Set up this model and workflow. Create a regular grid for `penalty` and `mixture` with 10 levels each; `mixture` should range from 0 to 1. For this assignment, we'll let `penalty` range from -5 to 5 (it's log-scaled).

How many total models will you be fitting when you fit these models to your folded data?

```
pokemon_spec <- multinom_reg(mixture = tune(), penalty = tune()) %>%
  set_mode("classification") %>%
  set_engine("glmnet")

pokemon_workflow <- workflow() %>%
  add_recipe(pokemon_recipe) %>%
  add_model(pokemon_spec)

penalty_grid <- grid_regular(penalty(range = c(-5, 5)),
                             mixture(range = c(0, 1)),
                             levels = 10)

penalty_grid
```

```
## # A tibble: 100 x 2
##       penalty mixture
##       <dbl>   <dbl>
## 1      0.00001      0
## 2      0.000129     0
## 3      0.00167      0
## 4      0.0215       0
## 5      0.278        0
## 6      3.59         0
## 7     46.4          0
## 8     599.          0
## 9    7743.          0
## 10 100000           0
## # ... with 90 more rows
```

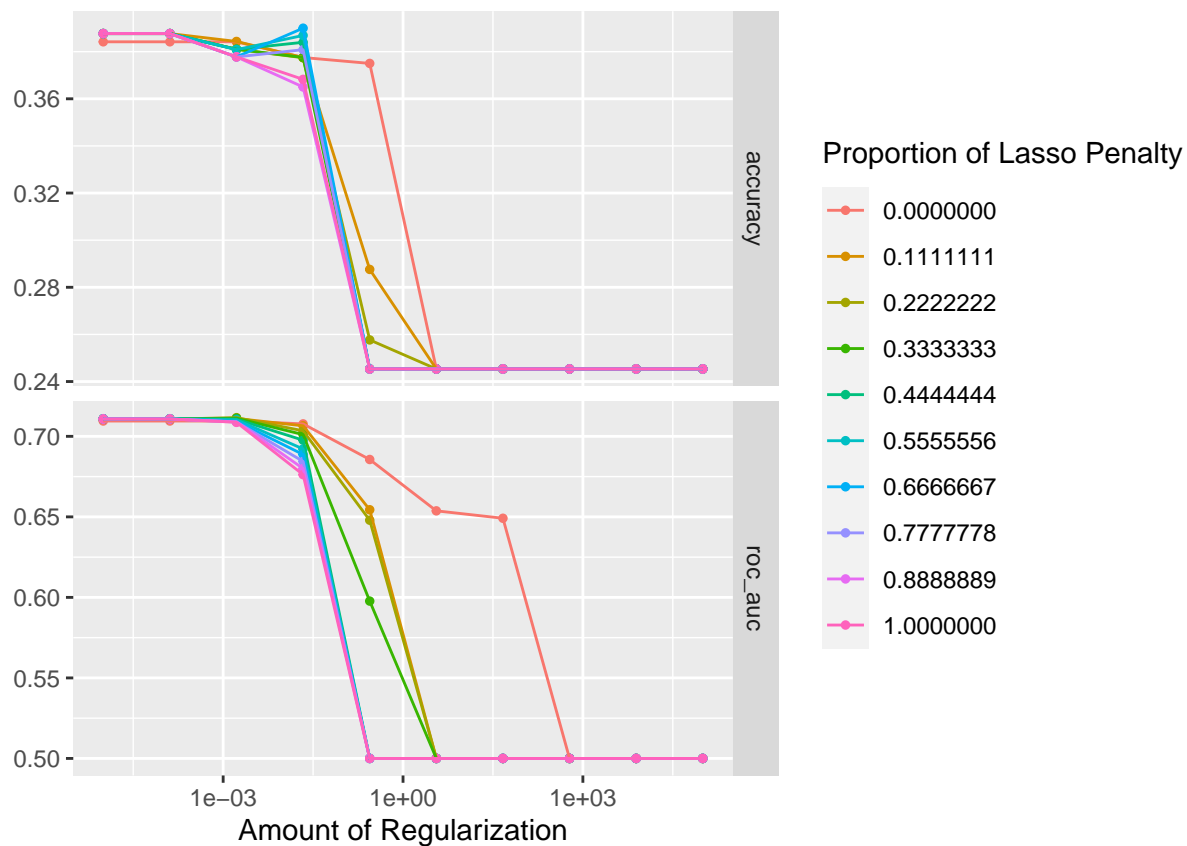
*There will be 500 models in total.*

## Exercise 6

Fit the models to your folded data using `tune_grid()`.

Use `autoplot()` on the results. What do you notice? Do larger or smaller values of `penalty` and `mixture` produce better accuracy and ROC AUC?

```
tune_res <- tune_grid(  
  pokemon_workflow,  
  resamples = pokemon_fold,  
  grid = penalty_grid  
)  
  
autoplot(tune_res)
```



*I notice that less lasso penalty produce higher accuracy and roc\_auc. And smaller mixture produce better accuracy and ROC AUC.*

## Exercise 7

Use `select_best()` to choose the model that has the optimal `roc_auc`. Then use `finalize_workflow()`, `fit()`, and `augment()` to fit the model to the training set and evaluate its performance on the testing set.

```
best_model <- select_best(tune_res, metric = "roc_auc")  
pokemon_final <- finalize_workflow(pokemon_workflow, best_model)  
pokemon_final_fit <- fit(pokemon_final, data = pokemon_train)
```

```
predicted_data <- augment(pokemon_final_fit, new_data = pokemon_test) %>%
  select(type_1, starts_with(".pred"))
```

## Exercise 8

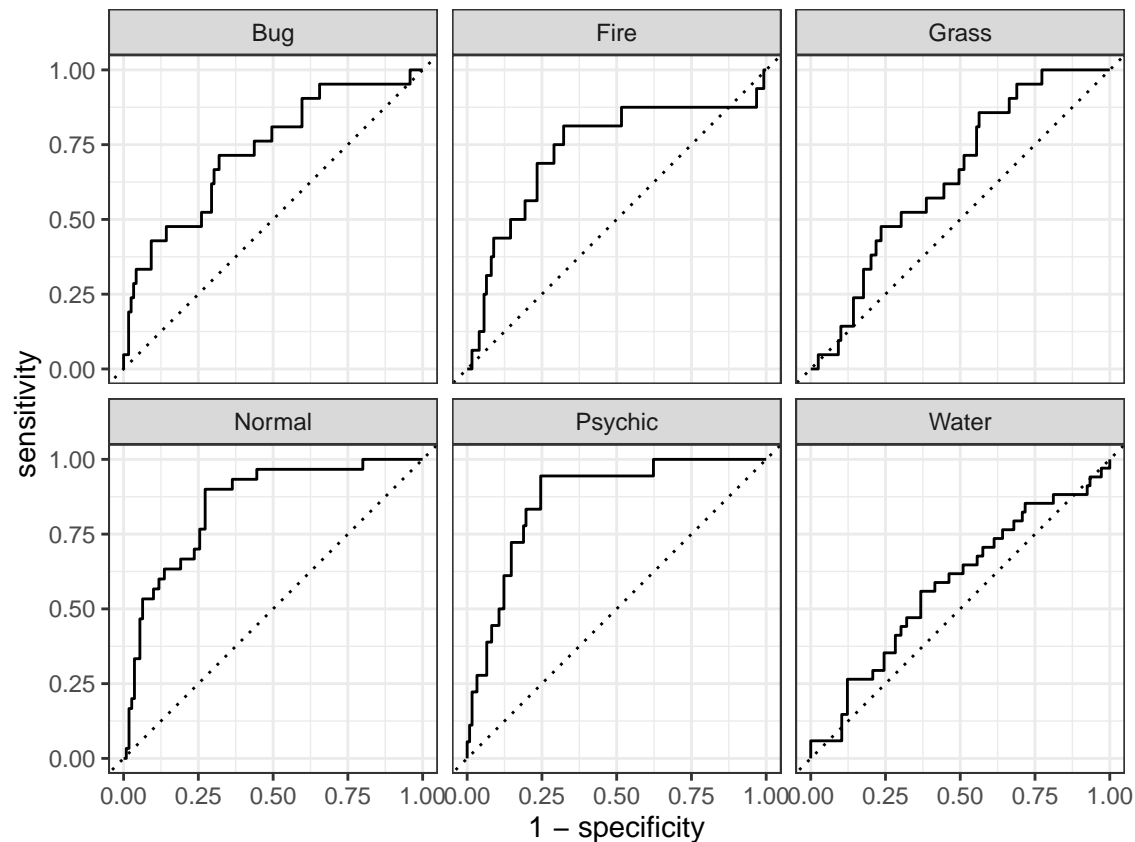
Calculate the overall ROC AUC on the testing set.

```
predicted_data %>% roc_auc(type_1, .pred_Bug:.pred_Water)
```

```
## # A tibble: 1 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>      <dbl>
## 1 roc_auc hand_till    0.728
```

Then create plots of the different ROC curves, one per level of the outcome. Also make a heat map of the confusion matrix.

```
predicted_data %>% roc_curve(type_1, .pred_Bug:.pred_Water) %>%
  autoplot()
```



```
augment(pokemon_final_fit, new_data = pokemon_test) %>%
  conf_mat(truth = type_1, estimate = .pred_class) %>%
  autoplot("heatmap")
```

Prediction	Bug -	9	0	7	4	0	5
	Fire -	0	1	0	0	0	1
	Grass -	0	1	2	0	3	0
	Normal -	5	2	0	19	1	9
	Psychic -	3	1	3	1	7	3
	Water -	4	11	9	6	7	16
		Bug	Fire	Grass	Normal	Psychic	Water
		Truth					

What do you notice? How did your model do? Which Pokemon types is the model best at predicting, and which is it worst at? Do you have any ideas why this might be?

*The overall roc\_auc is 0.70 which is not good enough. The model is best at predicting Normal type and water is the worst. This might be due to the resample techniques.*